



riminder

# INTRODUCTION TO DEEP LEARNING

---

**Mouhidine SEIV**  
FOUNDER - RIMINDER


Deep LEARNING A PRACTICAL COURSE  
ECOLE POLYTECHNIQUE, 05/04/2018

# History of Learning Systems

Évolution de l'intérêt pour cette recherche

Google Trends

- Apprentissage supervisé
- Apprentissage non supervisé
- Apprentissage par renforcement



Dans tous les pays. 01/01/2011 – 29/01/2018. Recherche sur le Web.

# Program & Course Logistics

---

- **Course 1 :** (05-04-18)
  - Introduction to Deep Learning - Mouhidine SEIV (Riminder)
- **Course 2 :** (12-04-18)
  - Deep Learning in Computer Vision - Slim FRIKHA (Riminder)
- **Course 3 :** (19-04-18)
  - Deep Learning in NLP - Paul COURSAUX (Riminder)
- **Course 4 :** (26-04-18)
  - Efficient Methods and Compression for Deep Learning - INVITED GUEST
- **Course 5:** (03-05-18)
  - Introduction to Deep Learning Frameworks - INVITED GUEST
- **Course 6:** (10-05-18)
  - Deployment in Production and Parallel Computing - INVITED GUEST

Location: Ecole Polytechnique from 6:30 pm to 7:30pm



<https://github.com/riminder>

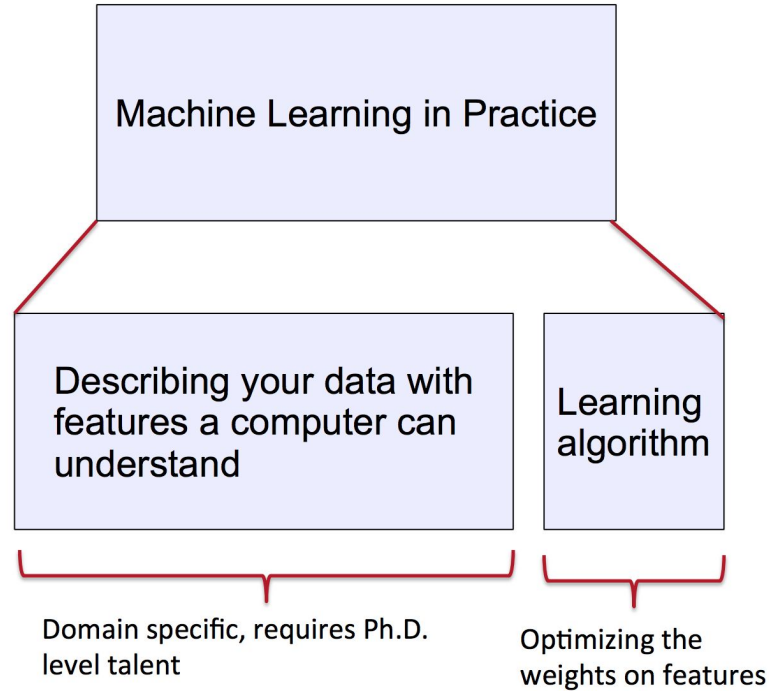
# Prerequisites

---

- Proficiency in python
- Linear algebra
- Basic probability and statistics
- Basic machine learning ( cost functions, derivatives , gradient methods)

# Machine Learning vs Deep Learning

---



# Reasons for Exploring Deep Learning


---

- In 2006 deep learning techniques started outperforming other Machine learning techniques. Why now?
- DL techniques benefit more from a lot of data
- Faster machines and multicore CPU/GPU help DL
- New models, algorithms, ideas

→ **Improved performance (first in speech and vision, then NLP)**


# What is Deep Learning?

- Representation learning attempts to automatically learn good features or representations
- Deep learning algorithms attempt to learn (multiple levels of) representation and an outputs
- From “raw” inputs  $\mathbf{x}$  (e.g. words)




# Deep Representations

- A **deep representation** is a composition of many functions



- Its gradient can be **backpropagated** by the chain rule





# Deep Neural Network

- A **deep neural network** is typically composed of:
  - Linear transformations

$$h_{k+1} = Wh_k$$

- Non-linear activation functions

$$h_{k+2} = f(h_{k+1})$$

- A loss function on the output, e.g.
  - Mean-squared error

$$l = ||y^* - y||^2$$

- Log likelihood


$$l = \log \mathbb{P}[y^*]$$

## Non-linear activation functions


Rectified linear unit




Maxout



Carefully tuned sigmoid



LSTM





# Training Neural Networks

- Sample gradient of expected loss  $L(\mathbf{w}) = \mathbb{E}[l]$

$$\frac{\partial l}{\partial \mathbf{w}} \sim \mathbb{E} \left[ \frac{\partial l}{\partial \mathbf{w}} \right] = \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}}$$



- Adjust  $\mathbf{w}$  down the sampled gradient

$$\Delta \mathbf{w} \propto \frac{\partial l}{\partial \mathbf{w}}$$



# Weight Sharing

- Recurrent neural network  
shares weights between time-steps



- Convolutional neural network  
shares weights between local regions

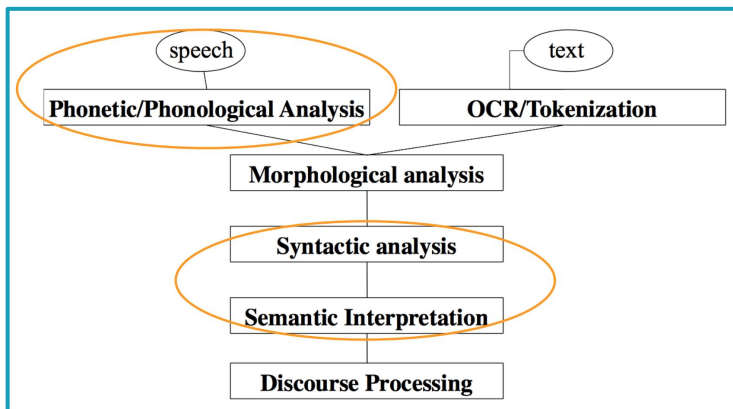
# What is NLP?

---


- Natural language processing is a field at the intersection of
  - Computer science
  - Artificial intelligence
  - and linguistics.
- Goal: for computers to process or “understand” natural language in order to perform tasks that are useful, e.g.
  - Question Answering
- Fully understanding and representing the meaning of language (or even defining it) is an illusive goal.
- Perfect language understanding is AI-complete



# Deep Learning for NLP




BEFORE Deep Learning




After Deep Learning

# What is Computer Vision?

- Most deep learning groups have (until 2 years ago) focused on computer vision
- Break through paper: ImageNet Classification with Deep Convolutional Neural Networks by Krizhevsky et al. 2012



This image is licensed under CC BY-NC-SA 2.0; changes made




This image is licensed under CC BY-SA 3.0; changes made



This image is licensed under CC BY-SA 2.0; changes made

Person  
Hammer





<https://github.com/riminder>



# Deep Learning for Computer Vision

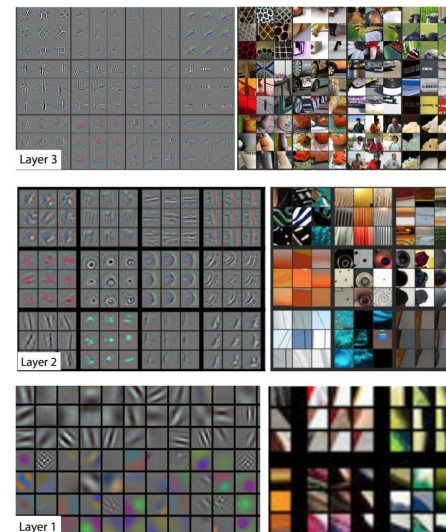
## BEFORE Deep Learning



Histogram of Gradients (HoG)  
Dalal & Triggs, 2005

Deformable Part Model  
Felzenswalb, McAllester, Ramanan, 2009

## After Deep Learning



Zeiler and Fergus (2013)




<https://github.com/riminder>

# Deep Learning for Speech

- The first breakthrough results of “deep learning” on large datasets happened in speech recognition
- Context-Dependent Pre-trained Deep Neural Networks for Large Vocabulary Speech Recognition Dahl et al. (2010)


Acoustic model	Recog \ WER	RT03S FSH	Hub5 SWB
Traditional features	1-pass -adapt	<b>27.4</b>	<b>23.6</b>
Deep Learning	1-pass -adapt	<b>18.5</b> (-33%)	<b>16.1</b> (-32%)





# Visualization

- Singular Value Decomposition (SVD)
- Principal component analysis (PCA)



$$\begin{array}{ccccc}
 \begin{array}{c} m \\ \boxed{\phantom{X}} \\ n \end{array} & & \begin{array}{c} r \\ \boxed{\begin{array}{c} | \\ U_1 \\ | \\ U_2 \\ | \\ U_3 \\ | \\ \vdots \end{array}} \end{array} & \begin{array}{c} r \\ \boxed{\begin{array}{ccc} S_1 & S_2 & 0 \\ & S_3 & \ddots \\ 0 & & S_r \end{array}} \end{array} & \begin{array}{c} m \\ \boxed{\begin{array}{c} \hline V_1 \hline \hline V_2 \hline \hline V_3 \hline \hline \vdots \hline \end{array}} \\ V^T \end{array} \\
X & = & U & S & \\
\begin{array}{c} m \\ \boxed{\phantom{\hat{X}}} \\ n \end{array} & & \begin{array}{c} k \\ \boxed{\begin{array}{c} | \\ U_1 \\ | \\ U_2 \\ | \\ U_3 \\ | \\ \vdots \end{array}} \end{array} & \begin{array}{c} k \\ \boxed{\begin{array}{ccc} S_1 & S_2 & 0 \\ & S_3 & \ddots \\ 0 & & S_k \end{array}} \end{array} & \begin{array}{c} m \\ \boxed{\begin{array}{c} \hline V_1 \hline \hline V_2 \hline \hline V_3 \hline \hline \vdots \hline \end{array}} \\ \hat{V}^T \end{array} \\
\hat{X} & = & \hat{U} & \hat{S} & 
\end{array}$$

$\hat{X}$  is the best rank  $k$  approximation to  $X$ , in terms of least squares.

- t-distributed stochastic neighbor embedding (T-SNE)  
<https://distill.pub/2016/misread-tsne/>


# Singular Value Decomposition (SVD)

## Corpus:

- I like deep learning. I like NLP. I enjoy flying.

```
import numpy as np
la = np.linalg
words = ["I", "like", "enjoy",
         "deep", "learnig", "NLP", "flying", "."]
X = np.array([[0, 2, 1, 0, 0, 0, 0, 0],
              [2, 0, 0, 1, 0, 1, 0, 0],
              [1, 0, 0, 0, 0, 0, 1, 0],
              [0, 1, 0, 0, 1, 0, 0, 0],
              [0, 0, 0, 1, 0, 0, 0, 1],
              [0, 1, 0, 0, 0, 0, 0, 1],
              [0, 0, 1, 0, 0, 0, 0, 1],
              [0, 0, 0, 0, 1, 1, 1, 0]])
```

```
U, s, Vh = la.svd(X, full_matrices=False)
```



- Printing first two columns of U corresponding to the 2 biggest singular values.

# Trouble with Learning systems

---



Adversarial example  
98% Toaster



Bias example  
95% Unqualified

# GENDER STEREOTYPING : WORD EMBEDDINGS

---

## HE

homemaker

nurse

receptionist

librarian

## SOFTBALL

pitcher

bookkeeper

registered nurse

waitress

## SHE

maestro

skipper

protege

philosopher

## FOOTBALL

footballer

businessman

maestro

cleric

# GENDER STEREOTYPING : REPRESENTATION

---

## Formulation:

- Let  $B$  be the gender subspace
- “Hard bias correction”: For a gender neutral word  $w$ , set  
 $w \leftarrow w - w_B / ||w - w_B||$
- “Soft Bias correction”: For word matrix  $W$ , gender neutral words  $N$  find a linear transformation  $T$

$$\min_T ||(TW)^T TW - W^T W||_F^2 + \lambda ||(TN)^T TB - ||_F^2$$

## Results:

- Reduced gender stereotyping in the word embeddings
- Performance on downstream tasks still almost the same

Man Is To Computer Programmer as Woman is to Homemaker? Debiasing Word Embeddings. T Bolukbasi, K-W Chang, J Zou, V Saligrama, A Kalai. Arxiv 2016

# Adversarial examples

(a) Image




(b) Prediction



(c) Adversarial Example




(d) Prediction




Original image

Temple (97%)



Perturbations



Adversarial example

Ostrich (98%)

# Generating Adversarial Examples

The Fast Gradient Sign Method

$$J(\tilde{\mathbf{x}}, \boldsymbol{\theta}) \approx J(\mathbf{x}, \boldsymbol{\theta}) + (\tilde{\mathbf{x}} - \mathbf{x})^\top \nabla_{\mathbf{x}} J(\mathbf{x}).$$

Maximize

$$J(\mathbf{x}, \boldsymbol{\theta}) + (\tilde{\mathbf{x}} - \mathbf{x})^\top \nabla_{\mathbf{x}} J(\mathbf{x})$$

subject to

$$\|\tilde{\mathbf{x}} - \mathbf{x}\|_\infty \leq \epsilon$$

$$\Rightarrow \tilde{\mathbf{x}} = \mathbf{x} + \epsilon \text{sign}(\nabla_{\mathbf{x}} J(\mathbf{x})).$$

Modern deep nets are very piecewise linear


Rectified linear unit




Maxout



Carefully tuned sigmoid




LSTM



# POINT OF CONTACT

---



Mouhidine SEIV  
CEO & Founder @Riminder  
[mouhidine.seiv@riminder.net](mailto:mouhidine.seiv@riminder.net)